

Assignment4

10.5 Exercise

1 How can you tell if an object is a tibble?

(Hint: try printing `mtcars`, which is a regular data frame).

```
mtcars

##           mpg cyl  disp  hp drat   wt  qsec vs am gear carb
## Mazda RX4      21.0   6  160.0  110 3.90 2.620 16.46 0  1    4    4
## Mazda RX4 Wag  21.0   6  160.0  110 3.90 2.875 17.02 0  1    4    4
## Datsun 710      22.8   4  108.0   93 3.85 2.320 18.61 1  1    4    1
## Hornet 4 Drive  21.4   6  258.0  110 3.08 3.215 19.44 1  0    3    1
## Hornet Sportabout 18.7   8  360.0  175 3.15 3.440 17.02 0  0    3    2
## Valiant         18.1   6  225.0  105 2.76 3.460 20.22 1  0    3    1
## Duster 360      14.3   8  360.0  245 3.21 3.570 15.84 0  0    3    4
## Merc 240D       24.4   4  146.7   62 3.69 3.190 20.00 1  0    4    2
## Merc 230        22.8   4  140.8   95 3.92 3.150 22.90 1  0    4    2
## Merc 280        19.2   6  167.6  123 3.92 3.440 18.30 1  0    4    4
## Merc 280C       17.8   6  167.6  123 3.92 3.440 18.90 1  0    4    4
## Merc 450SE      16.4   8  275.8  180 3.07 4.070 17.40 0  0    3    3
## Merc 450SL      17.3   8  275.8  180 3.07 3.730 17.60 0  0    3    3
## Merc 450SLC     15.2   8  275.8  180 3.07 3.780 18.00 0  0    3    3
## Cadillac Fleetwood 10.4   8  472.0  205 2.93 5.250 17.98 0  0    3    4
## Lincoln Continental 10.4   8  460.0  215 3.00 5.424 17.82 0  0    3    4
## Chrysler Imperial 14.7   8  440.0  230 3.23 5.345 17.42 0  0    3    4
## Fiat 128        32.4   4   78.7   66 4.08 2.200 19.47 1  1    4    1
## Honda Civic     30.4   4   75.7   52 4.93 1.615 18.52 1  1    4    2
## Toyota Corolla  33.9   4   71.1   65 4.22 1.835 19.90 1  1    4    1
## Toyota Corona   21.5   4  120.1   97 3.70 2.465 20.01 1  0    3    1
## Dodge Challenger 15.5   8  318.0  150 2.76 3.520 16.87 0  0    3    2
## AMC Javelin     15.2   8  304.0  150 3.15 3.435 17.30 0  0    3    2
## Camaro Z28      13.3   8  350.0  245 3.73 3.840 15.41 0  0    3    4
## Pontiac Firebird 19.2   8  400.0  175 3.08 3.845 17.05 0  0    3    2
## Fiat X1-9       27.3   4   79.0   66 4.08 1.935 18.90 1  1    4    1
## Porsche 914-2   26.0   4  120.3   91 4.43 2.140 16.70 0  1    5    2
## Lotus Europa    30.4   4   95.1  113 3.77 1.513 16.90 1  1    5    2
## Ford Pantera L  15.8   8  351.0  264 4.22 3.170 14.50 0  1    5    4
## Ferrari Dino    19.7   6  145.0  175 3.62 2.770 15.50 0  1    5    6
## Maserati Bora   15.0   8  301.0  335 3.54 3.570 14.60 0  1    5    8
## Volvo 142E      21.4   4  121.0  109 4.11 2.780 18.60 1  1    4    2

class(mtcars)

## [1] "data.frame"

class(as_tibble(mtcars))

## [1] "tbl_df"      "tbl"        "data.frame"
```

2 Compare and contrast the following operations on a data.frame and equivalent tibble. What is different? Why might the default data frame behaviours cause you frustration?

```
df <- data.frame(abc = 1, xyz = "a") df$xyz df[, "xyz"] df[, c("abc", "xyz")]
```

```
#data.frame
```

```
df <- data.frame(abc = 1, xyz = "a")
```

```
#tibble
```

```
tbl <- as.tibble(df)
```

```
df$xyz
```

```
## [1] a
```

```
## Levels: a
```

```
tbl$xyz
```

```
## Warning: Unknown or uninitialised column: 'xyz'.
```

```
## NULL
```

```
df[, "xyz"]
```

```
## [1] a
```

```
## Levels: a
```

```
tbl[, "xyz"]
```

```
## # A tibble: 1 x 1
```

```
##   xyz
```

```
##   <fct>
```

```
## 1 a
```

```
df[, c("abc", "xyz")]
```

```
##   abc xyz
```

```
## 1   1   a
```

```
tbl[, c("abc", "xyz")]
```

```
## # A tibble: 1 x 2
```

```
##   abc xyz
```

```
##   <dbl> <fct>
```

```
## 1  1.00 a
```

- Difference 1 Using \$ a data.frame will partially complete the column. So even though we wrote `df$xyz` it returned `df$xyz`. The advantage is it sometimes can save a few keystrokes, but on the other side, it can result in accidentally using a different variable than you thought you were using.
- Difference 2 With data.frames, with [the type of object that is returned differs on the number of columns. If it is one column, it won't return a data.frame, but instead will return a vector. With more than one column, then it will return a data.frame.

3 If you have the name of a variable stored in an object, e.g. `var <- "mpg"`, how can you extract the reference variable from a tibble?

```
df[[var]]
```

4 Practice referring to non-syntactic names in the following data frame by:

```
annoying <- tibble(  
  `1` = 1:10,  
  `2` = `1` * 2 + rnorm(length(`1`))  
)
```

```
annoying
```

```
## # A tibble: 10 x 2  
##       `1`     `2`  
##   <int> <dbl>  
## 1     1  2.99  
## 2     2  3.27  
## 3     3  6.54  
## 4     4  8.90  
## 5     5  9.85  
## 6     6 12.5  
## 7     7 14.9  
## 8     8 15.8  
## 9     9 16.9  
## 10    10 19.0
```

#(1) Extracting the variable called 1.

```
annoying[1]
```

```
## # A tibble: 10 x 1  
##       `1`  
##   <int>  
## 1     1  
## 2     2  
## 3     3  
## 4     4  
## 5     5  
## 6     6  
## 7     7  
## 8     8  
## 9     9  
## 10    10
```

```
annoying[["1"]]
```

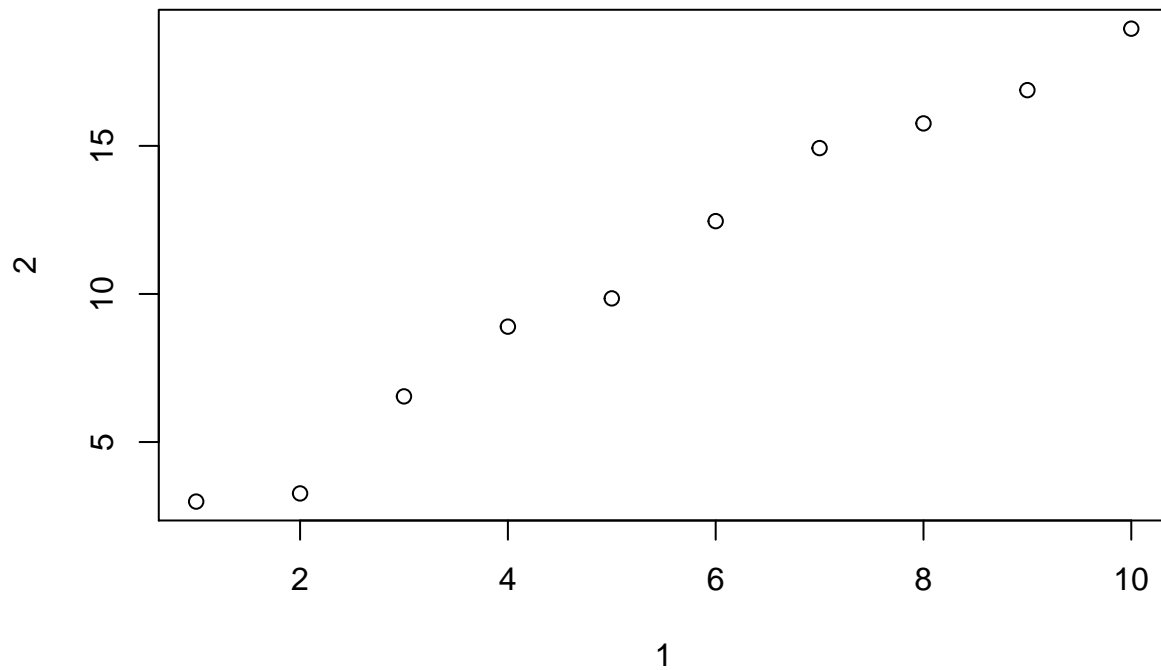
```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
annoying$'1'
```

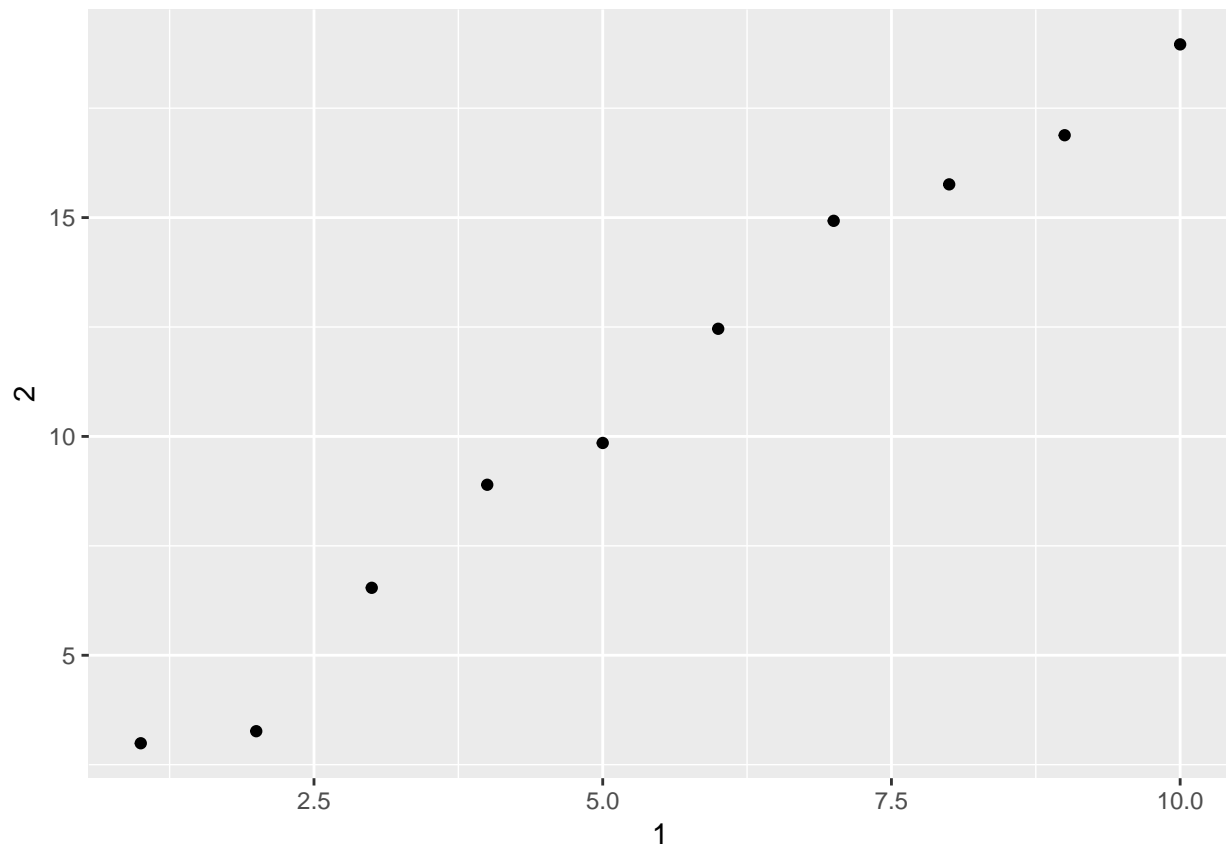
```
## [1] 1 2 3 4 5 6 7 8 9 10
```

#(2) Plotting a scatterplot of 1 vs 2.

```
plot(annoying)
```



```
ggplot(annoying, aes(x = `1`, y = `2`)) +geom_point()
```



```
#(3)Creating a new column called 3 which is 2 divided by 1.
annoying$`3` <- annoying$`2` / annoying$`1`
annoying[["3"]] <- annoying[["2"]] / annoying[["1"]]
```

```
#(4)Renaming the columns to one, two and three.
annoying <- rename(annoying, one = `1`, two = `2`, three = `3`)
glimpse(annoying)

## Observations: 10
## Variables: 3
## $ one    <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10
## $ two    <dbl> 2.989554, 3.265053, 6.541403, 8.895882, 9.850299, 12.459...
## $ three  <dbl> 2.989554, 1.632526, 2.180468, 2.223970, 1.970060, 2.0765...
```

5 What does `tibble::enframe()` do? When might you use it?

```
?enframe()
enframe(c(a = 5, b = 7, c=25))
```

```
## # A tibble: 3 x 2
##   name value
##   <chr> <dbl>
## 1 a     5.00
## 2 b     7.00
## 3 c    25.0
```

`enframe()` converts named atomic vectors or lists to two-column data frames. For unnamed vectors, the natural sequence is used as name column.

6 What option controls how many additional column names are printed at the footer of a tibble?

```
?print.tbl_df()
```

The print function for tibbles is in `print.tbl_df`: `n_extra` : Number of extra columns to print abbreviated information for, if the width is too small for the entire tibble. If `NULL`, the default, will print information about at most `tibble.max_extra_cols` extra columns

12.6.1 Exercise

```
who
```

```
## # A tibble: 7,240 x 60
##   country    iso2 iso3  year new_sp_m014 new_sp_m1524 new_sp_m2534
##   <chr>      <chr> <chr> <int>      <int>          <int>          <int>
## 1 Afghanistan AF    AFG   1980         NA            NA            NA
## 2 Afghanistan AF    AFG   1981         NA            NA            NA
## 3 Afghanistan AF    AFG   1982         NA            NA            NA
## 4 Afghanistan AF    AFG   1983         NA            NA            NA
## 5 Afghanistan AF    AFG   1984         NA            NA            NA
## 6 Afghanistan AF    AFG   1985         NA            NA            NA
## 7 Afghanistan AF    AFG   1986         NA            NA            NA
## 8 Afghanistan AF    AFG   1987         NA            NA            NA
```

```
## 9 Afghanistan AF AFG 1988 NA NA NA
## 10 Afghanistan AF AFG 1989 NA NA NA
## # ... with 7,230 more rows, and 53 more variables: new_sp_m3544 <int>,
## # new_sp_m4554 <int>, new_sp_m5564 <int>, new_sp_m65 <int>,
## # new_sp_f014 <int>, new_sp_f1524 <int>, new_sp_f2534 <int>,
## # new_sp_f3544 <int>, new_sp_f4554 <int>, new_sp_f5564 <int>,
## # new_sp_f65 <int>, new_sn_m014 <int>, new_sn_m1524 <int>,
## # new_sn_m2534 <int>, new_sn_m3544 <int>, new_sn_m4554 <int>,
## # new_sn_m5564 <int>, new_sn_m65 <int>, new_sn_f014 <int>,
## # new_sn_f1524 <int>, new_sn_f2534 <int>, new_sn_f3544 <int>,
## # new_sn_f4554 <int>, new_sn_f5564 <int>, new_sn_f65 <int>,
## # new_ep_m014 <int>, new_ep_m1524 <int>, new_ep_m2534 <int>,
## # new_ep_m3544 <int>, new_ep_m4554 <int>, new_ep_m5564 <int>,
## # new_ep_m65 <int>, new_ep_f014 <int>, new_ep_f1524 <int>,
## # new_ep_f2534 <int>, new_ep_f3544 <int>, new_ep_f4554 <int>,
## # new_ep_f5564 <int>, new_ep_f65 <int>, newrel_m014 <int>,
## # newrel_m1524 <int>, newrel_m2534 <int>, newrel_m3544 <int>,
## # newrel_m4554 <int>, newrel_m5564 <int>, newrel_m65 <int>,
## # newrel_f014 <int>, newrel_f1524 <int>, newrel_f2534 <int>,
## # newrel_f3544 <int>, newrel_f4554 <int>, newrel_f5564 <int>,
## # newrel_f65 <int>
```

```
who1 <- who %>%
  gather(new_sp_m014:newrel_f65, key = "key", value = "cases", na.rm = TRUE)
glimpse(who1)
```

```
## Observations: 76,046
## Variables: 6
## $ country <chr> "Afghanistan", "Afghanistan", "Afghanistan", "Afghanis..."
## $ iso2 <chr> "AF", "AF", "AF", "AF", "AF", "AF", "AF", "AF", "AF", ...
## $ iso3 <chr> "AFG", "AFG", "AFG", "AFG", "AFG", "AFG", "AFG", "AFG", "AFG"...
## $ year <int> 1997, 1998, 1999, 2000, 2001, 2002, 2003, 2004, 2005, ...
## $ key <chr> "new_sp_m014", "new_sp_m014", "new_sp_m014", "new_sp_m..."
## $ cases <int> 0, 30, 8, 52, 129, 90, 127, 139, 151, 193, 186, 187, 2...
```

```
who1

## # A tibble: 76,046 x 6
##   country iso2 iso3 year key cases
## * <chr> <chr> <chr> <int> <chr> <int>
## 1 Afghanistan AF AFG 1997 new_sp_m014 0
## 2 Afghanistan AF AFG 1998 new_sp_m014 30
## 3 Afghanistan AF AFG 1999 new_sp_m014 8
## 4 Afghanistan AF AFG 2000 new_sp_m014 52
## 5 Afghanistan AF AFG 2001 new_sp_m014 129
## 6 Afghanistan AF AFG 2002 new_sp_m014 90
## 7 Afghanistan AF AFG 2003 new_sp_m014 127
## 8 Afghanistan AF AFG 2004 new_sp_m014 139
## 9 Afghanistan AF AFG 2005 new_sp_m014 151
## 10 Afghanistan AF AFG 2006 new_sp_m014 193
## # ... with 76,036 more rows
```

```
who2 <- who1 %>%
  mutate(key = stringr::str_replace(key, "newrel", "new_rel"))
who2
```

```
## # A tibble: 76,046 x 6
```

```
##   country    iso2 iso3  year key      cases
##   <chr>      <chr> <chr> <int> <chr>    <int>
## 1 Afghanistan AF    AFG   1997 new_sp_m014    0
## 2 Afghanistan AF    AFG   1998 new_sp_m014   30
## 3 Afghanistan AF    AFG   1999 new_sp_m014    8
## 4 Afghanistan AF    AFG   2000 new_sp_m014   52
## 5 Afghanistan AF    AFG   2001 new_sp_m014  129
## 6 Afghanistan AF    AFG   2002 new_sp_m014   90
## 7 Afghanistan AF    AFG   2003 new_sp_m014  127
## 8 Afghanistan AF    AFG   2004 new_sp_m014  139
## 9 Afghanistan AF    AFG   2005 new_sp_m014  151
## 10 Afghanistan AF    AFG   2006 new_sp_m014  193
## # ... with 76,036 more rows
```

```
who3 <- who2 %>%
  separate(key, c("new", "type", "sexage"), sep = "_")
who3
```

```
## # A tibble: 76,046 x 8
##   country    iso2 iso3  year new  type  sexage cases
##   <chr>      <chr> <chr> <int> <chr> <chr> <chr>    <int>
## 1 Afghanistan AF    AFG   1997 new  sp    m014     0
## 2 Afghanistan AF    AFG   1998 new  sp    m014    30
## 3 Afghanistan AF    AFG   1999 new  sp    m014     8
## 4 Afghanistan AF    AFG   2000 new  sp    m014    52
## 5 Afghanistan AF    AFG   2001 new  sp    m014   129
## 6 Afghanistan AF    AFG   2002 new  sp    m014    90
## 7 Afghanistan AF    AFG   2003 new  sp    m014   127
## 8 Afghanistan AF    AFG   2004 new  sp    m014   139
## 9 Afghanistan AF    AFG   2005 new  sp    m014   151
## 10 Afghanistan AF    AFG   2006 new  sp    m014   193
## # ... with 76,036 more rows
```

```
who3 %>%
  count(new)
```

```
## # A tibble: 1 x 2
##   new      n
##   <chr> <int>
## 1 new  76046
```

```
who4 <- who3 %>%
  select(-new, -iso2, -iso3)
```

```
who5 <- who4 %>%
  separate(sexage, c("sex", "age"), sep = 1)
who5
```

```
## # A tibble: 76,046 x 6
##   country    year type  sex  age  cases
##   <chr>      <int> <chr> <chr> <chr> <int>
## 1 Afghanistan 1997 sp    m    014     0
## 2 Afghanistan 1998 sp    m    014    30
## 3 Afghanistan 1999 sp    m    014     8
## 4 Afghanistan 2000 sp    m    014    52
## 5 Afghanistan 2001 sp    m    014   129
## 6 Afghanistan 2002 sp    m    014    90
```

```
## 7 Afghanistan 2003 sp m 014 127
## 8 Afghanistan 2004 sp m 014 139
## 9 Afghanistan 2005 sp m 014 151
## 10 Afghanistan 2006 sp m 014 193
## # ... with 76,036 more rows
```

1 In this case study I set `na.rm = TRUE` just to make it easier to check that we had the correct values. Is this reasonable? Think about how missing values are represented in this dataset. Are there implicit missing values? What's the difference between an NA and zero?

```
who1 %>%
  filter(cases == 0) %>%
  nrow()
```

```
## [1] 11080
```

Yes, I think it is reasonable. Because there are zero's in the data, which means they may explicitly be indicating no cases. As for those values with NA values, they can either be really non-missing or be treated as non-missing by the data collector. So if we treat explicitly and implicitly missing values the same, we won't lose any information by dropping them.

2 What happens if you neglect the `mutate()` step? (`mutate(key = stringr::str_replace(key, "newrel", "new_rel"))`)

```
who3a <- who1 %>%
  separate(key, c("new", "type", "sexage"), sep = "_")
```

```
## Warning: Expected 3 pieces. Missing pieces filled with `NA` in 2580 rows
## [73467, 73468, 73469, 73470, 73471, 73472, 73473, 73474, 73475, 73476,
## 73477, 73478, 73479, 73480, 73481, 73482, 73483, 73484, 73485, 73486, ...].
```

`separate` causes the warning message saying the values are too few

```
filter(who3a, new == "newrel") %>% head()
```

```
## # A tibble: 6 x 8
##   country    iso2 iso3  year new    type sexage cases
##   <chr>      <chr> <chr> <int> <chr>  <chr> <chr>  <int>
## 1 Afghanistan AF    AFG   2013 newrel m014  <NA>   1705
## 2 Albania    AL    ALB   2013 newrel m014  <NA>    14
## 3 Algeria    DZ    DZA   2013 newrel m014  <NA>    25
## 4 Andorra    AD    AND   2013 newrel m014  <NA>     0
## 5 Angola     AO    AGO   2013 newrel m014  <NA>   486
## 6 Anguilla   AI    AIA   2013 newrel m014  <NA>     0
```

if we check the rows for keys beginning with "newrel_", we see that `sexage` is messing, and `type = m014`.

3 I claimed that iso2 and iso3 were redundant with country. Confirm this claim.

```
select(who3, country, iso2, iso3) %>%
  distinct() %>%
  group_by(country) %>%
  filter(n() > 1)

## # A tibble: 0 x 3
## # Groups:   country [0]
## # ... with 3 variables: country <chr>, iso2 <chr>, iso3 <chr>
```

4 For each country, year, and sex compute the total number of cases of TB. Make an informative visualization of the data.

```
who5 %>%
  group_by(country, year, sex) %>%
  filter(year > 1995) %>%
  summarise(cases = sum(cases)) %>%
  unite(country_sex, country, sex, remove = FALSE) %>%
  ggplot(aes(x = year, y = cases, group = country_sex, colour = sex)) +
  geom_line()
```

